Comprehensive Engineering Framework for High-Volume Multilingual AI Conversational Agents in Public Service and Enterprise Infrastructure

The architectural paradigm of voice-based artificial intelligence has undergone a fundamental shift toward real-time, low-latency streaming systems. Designing an agent capable of managing high-volume phone conversations for public service centers, grievance redressal, and outreach requires a sophisticated synthesis of telecommunications engineering, computational linguistics, and rigorous regulatory compliance. The development of such a system is most efficiently managed through the distribution of labor into three distinct technical verticals, each staffed by two specialists, to ensure deep focus on the linguistic cognitive layer, the telephony transport layer, and the operational governance layer.[1] This framework provides a stepwise methodology for building a multilingual AI calling agent designed for the complexities of the Indian linguistic landscape and the strictures of global telecommunications standards.

Vertical One: Linguistic Intelligence and Cognitive Architecture

The primary responsibility of the first vertical is the development of the "brain" and "sensory" apparatus of the AI agent. This involves the orchestration of the Speech-to-Text (STT), Large Language Model (LLM), and Text-to-Speech (TTS) pipeline. The core technical objective for this team is the reduction of end-to-end latency to sub-500 milliseconds while maintaining high accuracy across diverse dialects and code-mixed speech.[1]

Real-Time Speech Recognition in Multilingual Environments

For a public service agent operating in India, the STT component must transcend standard monolingual recognition. The linguistic engineers must implement an engine capable of handling code-switching, particularly Hinglish, and regional variations that occur every 50 to 100 kilometers.[5] Global models frequently exhibit a "functional failure" in this domain, with error rates in languages like Tamil or Maithili exceeding 55%, compared to the 5-6% seen in standardized Hindi.[5] The team should prioritize indigenous models such as Sarvam AI's Saarika or Bhashini's ASR, which have been benchmarked to outperform global competitors by over 50 percentage points in Indian contexts.[5]

STT Model	Latency Profile	Indian Language Accuracy	Code-Mixing Support	Dialect Robustness
Deepgram Nova-3	< 100 ms	High (Hindi/English)	Moderate	Low
Sarvam Saarika-v2.5	Medium-Low	Exceptional (11+ Lang)	High	High
Bhashini ASR	Low (Streaming)	High (22 Lang)	Very High	Moderate
OpenAI Whisper	500ms - 2000ms	Low (Dravidian gap)	Low	Low

The technical implementation must utilize a streaming architecture rather than a cascading one. In a cascading pipeline, the system waits for the user to finish speaking before transcribing, which introduces cumulative delays. A streaming architecture processes audio in small chunks, transcribing incrementally so that the LLM can begin intent detection before the utterance is complete.[1] This is achieved through the use of WebSockets or gRPC streams to maintain a persistent bi-directional connection between the caller and the transcription server.[8]

Cognitive Memory Systems and RAG Integration

For use cases such as grievance redressal or complex customer support, the agent requires a tiered memory architecture to maintain context across sessions. The engineers must implement a dual-layer memory system comprising thread-level (short-term) memory and cross-thread (long-term) memory.[10] Thread-level memory acts as the agent's working memory, managed through checkpoints in frameworks like LangGraph to ensure continuity within a single call.[10] Cross-thread memory utilizes Retrieval-Augmented Generation (RAG) to store citizen preferences, previous complaint history, and relevant policy documents in a vector database like ChromaDB.[11]

Memory Layer	Type	Persistence Mechanism	Primary Function
Short-Term	Thread-level	In-Memory/Checkpoints	Maintaining session context
Long-Term	Cross-thread	Vector DB (JSON/Semantic)	Retaining user history/preferences
Policy RAG	Static/Dynamic	Knowledge Base Sync	Grounding responses in regulations
Entity Layer	Structured	CRM/Database Integration	Extracting names, IDs, and dates

The implementation of a four-layer RAG system is recommended to solve the "context dilution" problem. This involves extracting summarized facts, entity recognition logs, and thematic clusters from historical interactions.[11] In a government grievance context, this allows the agent to recognize that a caller is following up on a previously reported water shortage, significantly reducing the citizen's effort and improving the first-call resolution rate.[12]

Neural Voice Synthesis and Emotional Prosody

The final stage of the linguistic pipeline is the conversion of text back into speech. To be effective in public service, the agent's voice must not sound "robotic," which is a common point of failure for legacy IVR systems. The team must utilize neural TTS models that support emotional expression and prosody, allowing the agent to convey empathy when a citizen is frustrated or urgency during a time-sensitive outreach campaign.[1]

TTS Provider	Time to First Byte (TTFB)	Voice Quality (MOS)	Emotional Range	Language Breadth

ElevenLabs Flash	< 100 ms	4.5+	Exceptional	30+
Cartesia Sonic 3	< 100 ms	4.4	Very Expressive	40+
Sarvam Bulbul v3	Medium-Low	4.2	Balanced	11+ (Indic)
Indic-Parler (OS)	Variable	4.0	Natural	23+ (Indic)

Latency optimization in the TTS layer is achieved by beginning audio playback while the complete response is still being generated. This "speaking as it thinks" approach is facilitated by streaming the LLM tokens directly into the TTS engine.[1] For Indian vernacular languages, the use of open-source models like Indic-Parler, released under the Apache 2.0 license, provides a cost-effective way to achieve realistic, expressive speech in 23 languages without the high per-minute costs of proprietary global platforms.[7]

Vertical Two: Telephony Infrastructure and Latency Optimization

The second vertical is tasked with building the "nervous system" of the agent, focusing on the high-volume transport of audio data and the orchestration of real-time events. This team manages the SIP (Session Initiation Protocol) trunking, the media gateway, and the interruption handling logic that allows for natural turn-taking in conversation.[2]

SIP Trunking and PSTN Connectivity

For high-volume inbound and outbound campaigns, carrier-grade SIP trunking is the non-negotiable standard for reliability and scalability.[2] Unlike experimental WebSocket-only deployments, SIP trunks connect the AI platform directly to the Public Switched Telephone Network (PSTN), enabling the agent to send and receive millions of calls with enterprise-grade Quality of Service (QoS).[2] The engineers must configure an Elastic SIP stack that can handle "instant bursts" during peak hours, such as during a product launch or an emergency public service announcement.[16]

Connectivity Type	Reliability	Scalability	Best For	Technical Constraint
SIP Trunk	Carrier-grade	100,000s calls	Production PSTN calls	Requires Port 5060/5061
WebRTC	High (App-based)	High (Browser)	In-app support	NAT/Firewall traversal
WebSocket	Experimental	Limited	Prototypes/Web widgets	Jitter sensitivity
WhatsApp Voice	Moderate	Growing	Social engagement	Meta API limitations

The telephony team must establish a media gateway—using Asterisk, FreePBX, or specialized cloud providers like Telnyx—to bridge the gap between traditional telecom networks and the real-time AI pipeline.[2] This gateway is responsible for transcoding audio into formats required by STT engines (e.g., 8kHz or 16kHz linear PCM) and managing the "Signaling" and "Media" planes separately to ensure low-latency transport.[2]

Advanced Interruption and Turn-Taking Logic

A critical failure point for AI voice agents is the "barge-in" problem, where the AI continues speaking after the user has interrupted it. This vertical must implement a robust Turn-Taking Model that mimics human listening behavior.[19] This requires the integration of high-performance Voice Activity Detection (VAD) at the edge of the telephony gateway.[19] The moment the VAD detects human speech, the system must trigger an immediate "stop" signal to the TTS synthesis and the audio playback buffer.[8]

Effective interruption handling is mathematically governed by the latency of the VAD signal. For an interaction to feel natural, the bot must stop talking within 200 milliseconds of the user beginning an interruption.[4] The team should also implement "listening" while "speaking" capabilities, where the STT engine continues to transcribe the user's background noise or verbal cues during the bot's speaking turn, allowing the LLM to adjust its response or

acknowledge the interruption gracefully.[8]

Orchestration and Latency Budgets

The orchestration layer is the software "conductor" that manages the flow of data between STT, LLM, and TTS components. To achieve sub-400ms latency, the team must treat the entire pipeline as a single coordinated system rather than three independent API calls.[4] This involves optimizing the "inter-token latency" (ITL) and the "Time to First Byte" (TTFB) across the entire stack.[1]

Latency is calculated as:

$$L_{total} = L_{transcription} + L_{inference} + L_{synthesis} + L_{network}$$

To meet the 500ms threshold, each component must operate within strict budgets:

- **STT Latency:** < 150 ms.[1]
- **LLM TTFT:** < 100 ms (achieved through KV-caching and model quantization).[4]
- **TTS TTFB:** < 100 ms.[1]
- **Network/Telephony Hops:** < 100 ms.[1]

The orchestration team should implement "Parallelism at the Edge," where the system begins generating the TTS audio for the first sentence of an LLM response while the second sentence is still being processed by the model.[1]

Vertical Three: Compliance, Governance, and Analytics

The third vertical is responsible for the ethical, legal, and operational integrity of the AI calling agent. This team focuses on regulatory adherence to TRAI and DPDPA, the design of human-in-the-loop (HITL) escalation protocols, and the generation of actionable insights through an analytics dashboard.[21]

Regulatory Frameworks and TRAI Compliance

For outbound calling agents in India, strict adherence to the Telecom Regulatory Authority of India (TRAI) guidelines is essential to prevent penalties and blacklisting. The team must ensure that all promotional calls use the 140-series number prefix, while all transactional and service-related calls use the 160-series.[3] The 160-series is specifically designated for government and financial institutions to build trust and ensure high answer rates.[24]

Number Series	Purpose	DND Enforcement	Eligibility
140 Series	Promotional/Marketing	Mandatory Scrubbing	Any registered telemarketer
160 Series	Service/Transactional	Bypassed for critical info	Banks, Govt, Insurers
10-Digit Mobile	Personal/Inbound	Not for outbound marketing	Individuals/Small Biz

Compliance engineers must integrate the system with the Distributed Ledger Technology (DLT) platform, a blockchain-based registry where all call headers and voice templates must be approved before deployment.[3] They must also implement real-time "Scrubbing" against the National Customer Preference Register (NCPR) to ensure that users on the Do Not Disturb (DND) list are not contacted without explicit, verifiable consent.[3]

Data Privacy and the DPDPA 2023

The agent must operate within the legal boundaries of the Digital Personal Data Protection Act (DPDPA) 2023. This mandates that all data processing must be based on "free, specific, informed, and unambiguous" consent.[21] Every call initiated by the AI must begin with a clear disclosure that the user is interacting with an AI bot.[3]

Furthermore, the vertical must implement technical safeguards, including encryption of voice logs and data masking, to protect personal information.[21] The DPDPA also grants "Data Principals" (the citizens) the right to access, correct, or erase their data.[21] The system must be architected to support "The Right to be Forgotten," allowing for the automated deletion of voice records and transcripts upon user request or after the fulfillment of the processing purpose.[21]

Human-in-the-Loop Escalation and Context Transfer

No AI agent can handle 100% of complex grievances or emotional distress. The team must design a "Seamless Handoff" protocol that transfers the call to a human specialist when specific triggers are met.[22] These triggers include low model confidence, detected sentiment of anger or frustration, or an explicit request for a "human".[22]

A critical component of this handoff is the Context Transfer Object. Rather than forcing the human agent to listen to the entire previous conversation, the AI must generate a structured summary in real-time.[34]

Field	Description	Importance
Intent Summary	Concise description of the user's primary goal	Reduces repetition
Sentiment Velocity	Measure of emotional shift (e.g., +0.2 to -0.5)	Prepares agent for mood
Data Collected	Extracted entities like Account ID or Address	Streamlines verification
Escalation Reason	Why the AI failed (e.g., "Complex Policy Query")	Identifies specific blocker
Recommended Action	Next step for the human to take	Accelerates resolution

The human agent’s dashboard should display this information through a "context card" via a WebSocket connection, ensuring it is visible before the call is connected.[35] This creates a "magical" experience where the human agent greets the citizen by name and acknowledges the specific issue already discussed.[22]

Analytics Generation and Performance KPIs

The success of the agent workforce must be measured through a specialized set of Key Performance Indicators (KPIs) tailored for conversational AI.[23] The analytics team must build a real-time dashboard that tracks these metrics to identify gaps in conversation design and system performance.

KPI	Formula / Calculation	Target Benchmark
Containment Rate	(Resolved by AI / Total Calls) x 100	> 65% [37]
First Call Resolution	% resolved in one interaction	70 - 80% [38]
Sentiment Velocity	Change in sentiment score / Time period	> -0.1 (Stable) [23]
Bot Accuracy (WER)	(Correct / Total Words) x 100	> 90% [38]
Response Time	End of User Speech to Start of AI Audio	< 500 ms [37]
Cost per Resolution	Total infrastructure cost / Resolved calls	Industry-dependent

The dashboard should include a "Performance Heatmap" showing hourly surges in call volume and system latency, allowing the infrastructure team to adjust GPU allocation or SIP capacity dynamically.[23] Advanced sentiment analysis models must monitor "Sentiment Velocity" every 30 seconds to catch deteriorating conversations before the user disconnects in frustration.[23]

Step-by-Step Implementation Process

The project is divided into four chronological phases, with tasks distributed across the three verticals to ensure parallel progress.

Phase 1: Infrastructure and Basic Pipeline (Weeks 1-4)

The goal of this phase is to establish the fundamental connectivity and the "Hello World" of the voice agent.

- **Vertical 1 (NLP):** Evaluate and select the STT and TTS engines for the primary target languages (e.g., Hindi, English, Tamil). Establish basic prompt logic for intent classification.[1]
- **Vertical 2 (Telephony):** Secure a SIP trunk and configure the media gateway. Set up a WebSocket server to ingest audio streams from the gateway to the STT engine.[2]
- **Vertical 3 (Compliance):** Register the "Principal Entity" on the DLT platform. Begin the application for the 160-series number block.[3]

Phase 2: Cognitive Depth and Contextual Memory (Weeks 5-10)

This phase focuses on making the agent intelligent and context-aware.

- **Vertical 1 (NLP):** Implement the four-layer RAG architecture. Connect the system to the public service policy database and set up cross-thread memory.[10]
- **Vertical 2 (Telephony):** Develop the turn-taking logic and interruption handling. Optimize the orchestration layer to achieve a sub-500ms latency "budget".[8]
- **Vertical 3 (Compliance):** Design the HITL escalation flows and the JSON context transfer schema. Finalize the AI disclosure scripts for legal compliance.[3]

Phase 3: Analytics and Operational Hardening (Weeks 11-16)

This phase prepares the system for high-volume production and performance monitoring.

- **Vertical 1 (NLP):** Fine-tune the LLM for domain-specific terminology (e.g., grievance categories). Implement code-mixing support for regional dialects.[5]
- **Vertical 2 (Telephony):** Conduct load testing to simulate 1,000+ concurrent calls. Optimize the auto-scaling triggers for the GPU inference cluster.[17]
- **Vertical 3 (Compliance):** Build the Analytics Dashboard and integrate with the CRM. Launch a limited pilot to gather baseline KPI data.[23]

Phase 4: Full-Scale Deployment and Optimization (Weeks 17-20)

The final phase involves the full rollout and the establishment of a continuous feedback loop.

- **Vertical 1 (NLP):** Analyze failed interactions to retrain intent recognition models. Update the knowledge base based on "Knowledge Gaps" identified during the pilot.[40]
- **Vertical 2 (Telephony):** Monitor SIP trunk uptime and call completion rates. Fine-tune the VAD sensitivity to handle diverse background noise environments.[16]
- **Vertical 3 (Compliance):** Conduct a formal DPDPA audit and ensure all voice logging meets the 1-year retention mandate. Transition to full 140/160 series outbound operations.[21]

Deep Insight: The Convergence of Sovereign AI and Local Telephony

A distinguishing feature of high-performance agents in non-Western markets is the reliance on "Sovereign Cloud" and "Licensed Telco" integration. For public service centers in India, data localization is not merely a preference but often a regulatory requirement.[27] Deploying on platforms like Bhashini or Sarvam AI, which are hosted within Indian data centers, ensures that sensitive citizen grievances regarding healthcare, taxes, or legal matters are processed without traversing international boundaries.[44]

Furthermore, the "Sovereign AI" approach allows for the development of "Dialect-First" models. Standard models often fail to capture the prosody of rural speech, leading to high abandonment rates in public service campaigns targeted at non-urban populations.[5] By leveraging the 15,000+ hours of transcribed data from Bhashini's nationwide initiative, the development team can fine-tune the agent for the specific linguistic nuances of over 400 districts.[46]

Technical Logic of High-Volume Outbound Campaigns

In an outreach campaign—such as a large-scale survey or a vaccination reminder—the system must manage "Contact Rates" alongside conversation quality. A high-volume dialer must be "Answering Machine Detection" (AMD) aware.

The mathematical model for outbound efficiency is defined by:

$$E_{outbound} = (C_{answered} / C_{dialed}) \times (R_{valid} / C_{answered})$$

Where:

- C_{dialed} is the total calls initiated.
- $C_{answered}$ is the calls where a person (not a machine) picked up.[48]
- R_{valid} is the number of interactions that completed the intended task (survey response, reminder confirmation).[37]

The AI agent must be trained to detect the characteristic "Beep" of an answering machine within 500ms and either hang up to save costs or leave a pre-synthesized personalized message.[17] In India, the use of the 160-series for these calls increases the $C_{answered}$ rate by up to 20-30%, as citizens are more likely to trust the verified numeric prefix.[19]

Conclusion: The Integrated Future of Voice Automation

The engineering of a high-volume AI calling agent is no longer a search for a better model but a search for a better system. By orchestrating a sub-400ms streaming pipeline with deep contextual memory and carrier-grade telephony, organizations can move beyond simple automated menus into the realm of natural, human-like dialogue.[1] For public service and enterprise centers, this technology represents a significant reduction in operational overhead—estimated at up to 60-80%—while simultaneously improving the accessibility of services to millions of multilingual users.[13] The methodology of dividing the project into linguistic, telephony, and compliance verticals ensures that the resultant system is not only

technologically advanced but also socially responsible and legally sound.[3]

Works cited

1. The voice AI stack for building agents in 2026 - AssemblyAI, accessed March 7, 2026, https://www.assemblyai.com/blog/the-voice-ai-stack-for-building-agents
2. Best Telephony Integration for AI Call Centers in 2026 ... - Subverse AI, accessed March 7, 2026, https://subverseai.com/blogs/best-telephony-integration-for-ai-call-centers-in-2026-sip-vs-web-call-vs-whatsapp
3. TRAI Updates for AI Calling: Your Complete Compliance Guide for 2026, accessed March 7, 2026, https://qcall.ai/trai-updates-for-ai-calling
4. Building Real-Time Voice AI: Inside the Infrastructure That Achieves ..., accessed March 7, 2026, https://simplismart.ai/blog/real-time-voice-ai-sub-400ms-latency
5. Global AI Firms Lag Indian Rivals In Speech Recognition, Benchmark Finds, accessed March 7, 2026, https://www.businessworld.in/article/global-ai-firms-lag-indian-rivals-in-speech-recognition-benchmark-finds-593680
6. Global speech AI struggles to understand India: Report - The Economic Times, accessed March 7, 2026, https://m.economictimes.com/small-biz/security-tech/technology/global-speech-ai-struggles-to-understand-india-report/articleshow/128410287.cms
7. Multilingual Voice AI in India: STT–LLM–TTS Pipeline ... - Subverse AI, accessed March 7, 2026, https://subverseai.com/blogs/multilingual-voice-ai-in-india-stt-llm-tts-pipeline-vs-speech-to-speech-s2s-what-should-you-choose
8. Architecting Low-Latency, Real-Time AI Voice Agents: Challenges & Solutions - Dev.to, accessed March 7, 2026, https://dev.to/lifeisverygood/architecting-low-latency-real-time-ai-voice-agents-challenges-solutions-hdn
9. Bhashini.ai API Docs, accessed March 7, 2026, https://www.bhashini.ai/docs
10. FareedKhan-dev/langgraph-long-memory: A detail Implementation of handling long-term memory in Agentic AI - GitHub, accessed March 7, 2026, https://github.com/FareedKhan-dev/langgraph-long-memory
11. Built a four-layer RAG memory system for my AI agents (solving the context dilution problem), accessed March 7, 2026, https://www.reddit.com/r/learnmachinelearning/comments/1rebes6/built_a_fourlayer_rag_memory_system_for_my_ai/
12. RAG-Based Multi-Agent Framework for Intelligent Government Grievance Redressal Using Small Language Models | Request PDF - ResearchGate, accessed March 7, 2026, https://www.researchgate.net/publication/400186314_RAG-Based_Multi-Agent_Framework_for_Intelligent_Government_Grievance_Redressal_Using_Small_Language_Models
13. RAG Agents: Revolutionizing Voice and IVR Systems - Teneo.Ai, accessed March 7,

2026, https://www.teneo.ai/blog/rag-agents-revolutionizing-voice-and-ivr-systems
14. Speech Synthesis - AI4Bharat, accessed March 7, 2026, https://ai4bharat.iitm.ac.in/areas/tts
15. Indic Parler TTS - AI4Bharat, accessed March 7, 2026, https://ai4bharat.iitm.ac.in/areas/model/TTS/Indic%20Parler%20TTS
16. What Is SIP Trunking? Benefits, How It Works, and More | Vonage, accessed March 7, 2026, https://www.vonage.com/resources/articles/what-is-sip-trunking-part-1-the-sip-definition/
17. Hire AI Voice Agents to Automate Your Phone Calls | UnleashX, accessed March 7, 2026, https://www.unleashx.ai/voice-ai/
18. VoiceInfra - Conversational AI Voice Agents for Enterprise | Voice AI Platform, accessed March 7, 2026, https://voiceinfra.ai/
19. I tested Bland, Vapi, and Retell AI for a month on live calls. Here is why I'm migrating everything to Retell AI. : r/ProductivityApps - Reddit, accessed March 7, 2026, https://www.reddit.com/r/ProductivityApps/comments/1qiwed9/i_tested_bland_vapi_and_retell_ai_for_a_month_on/
20. End-to-End Text-to-Speech: Cut Voice Latency by 50-70% - Deepgram, accessed March 7, 2026, https://deepgram.com/learn/unified-end-to-end-text-to-speech-architecture-cuts-latency
21. India's Digital Personal Data Protection Act (DPDPA) - TrustArc, accessed March 7, 2026, https://trustarc.com/resource/indias-digital-personal-data-protection-act-dpdpa/
22. Human in the Loop for Voice AI: Staying Safe Without Slowing Down - Nurix AI, accessed March 7, 2026, https://www.nurix.ai/blogs/human-in-the-loop-voice-ai
23. AI Contact Center KPIs: 7 Must-Track Voicebot Metrics Guide 2026 - QCall AI, accessed March 7, 2026, https://qcall.ai/ai-contact-center-kpis
24. TRAI 160 Series Deadline: How to Ensure Compliant, Trusted Voice Communication, accessed March 7, 2026, https://ozonetel.com/160-series-deadline-looms/
25. TRAI 160 Number Series : A Complete Guide for Businesses in 2026 - cloud communications solution, accessed March 7, 2026, https://www.acefone.com/blog/all-about-160-number-series/
26. TRAI 160-Series: Fintech's New Verified Calling Standard - TeleCMI, accessed March 7, 2026, https://www.telecmi.com/blog/trai-160-number-series-explained
27. India Outbound Call Regulations 2025: Complete DPDPA & TRAI Compliance Guide | TALK-Q, accessed March 7, 2026, https://talk-q.com/outbound-call-regulations-in-india
28. Trai Guidelines for Outbound Calling Timings in India - 2026 - ClearTouch, accessed March 7, 2026, https://www.cleartouch.in/blog/trai-guidelines-for-outbound-calling-timings-in-i

ndia/
29. THE DIGITAL PERSONAL DATA PROTECTION ACT, 2023 (NO. 22 OF 2023) An Act to provide for the processing of digital personal data in, accessed March 7, 2026, https://www.meity.gov.in/static/uploads/2024/06/2bf1f0e9f04e6fb4f8fef35e82c42aa5.pdf
30. Is AI Calling Legal in the U.S., EU, India & Beyond? [2025 Guide], accessed March 7, 2026, https://botpenguin.com/blogs/is-ai-calling-legal
31. India Passes the Digital Personal Data Protection Rules, Ushering in a New Digital Age in India | Privacy World, accessed March 7, 2026, https://www.privacyworld.blog/2025/11/india-passes-the-digital-personal-data-protection-rules-ushering-in-a-new-digital-age-in-india/
32. The Ultimate Guide to Chatbots for Customer Service [2025] | AI chatbot - Puzzel, accessed March 7, 2026, https://www.puzzel.com/blog/chatbots-guide
33. 7 best practices for human handoff in chat support (2025 guide) - eesel AI, accessed March 7, 2026, https://www.eesel.ai/blog/best-practices-for-human-handoff-in-chat-support
34. AI Escalation Management: Mastering Handoffs in Modern Support - SearchUnify, accessed March 7, 2026, https://www.searchunify.com/resource-center/sudo-technical-blogs/ai-escalation-management-mastering-handoffs-in-modern-support/
35. designing ai agent handoffs to humans, what's the least jarring approach - Reddit, accessed March 7, 2026, https://www.reddit.com/r/AI_Agents/comments/1r6r54e/designing_ai_agent_handoffs_to_humans_whats_the/
36. AI-Human Call Handoff Protocols: Engineering Seamless Transitions in Hybrid Systems, accessed March 7, 2026, https://smith.ai/blog/ai-human-call-handoff-protocols
37. KPI Chatbot Metrics: Essential Guide to Tracking AI Success in 2025 - Dialzara, accessed March 7, 2026, https://dialzara.com/blog/ai-chatbot-kpis-what-to-track-in-2025
38. KPIs for AI Voice Agents in Contact Centers | Key Metrics, accessed March 7, 2026, https://www.hakunamatatatech.com/our-resources/blog/ai-voice-agents-in-contact-centers
39. How to Measure KPIs for AI Customer Support: A Comprehensive Guide - Helply, accessed March 7, 2026, https://helply.com/blog/measure-kpi-for-ai-customer-support
40. Top KPIs to Track After Deploying a Voice AI Agent | Vozzo.AI Blog, accessed March 7, 2026, https://vozzo.ai/blog-top-kpis-to-track-after-deploying-a-voice-ai-agent
41. Why Your Business Should Move to 160-Series Numbers for Service Calls? - ClearTouch, accessed March 7, 2026, https://www.cleartouch.in/blog/why-business-should-move-to-160-series-numbers-for-service-calls/
42. Best Text to Speech AI Software for Multilingual Voices - Devnagri AI, accessed

March 7, 2026, https://devnagri.com/text-to-speech/
43. How AI Is Transforming B2B Customer Support | Complete Guide for 2025 - Pylon, accessed March 7, 2026, https://www.usepylon.com/blog/ai-transforming-b2b-customer-support-2025
44. Sarvam AI Scores Higher Than Gemini & OpenAI on Indian Document Reading Benchmarks, accessed March 7, 2026, https://thebetterindia.com/innovation/indian-ai-document-reading-sarvam-gemini-openai-language-tests-11092770
45. Empowering Startups to Build Innovative Language Solutions with Multilingual AI for Digital Inclusion - Bhashini, accessed March 7, 2026, https://bhashini.gov.in/udbhav
46. AI4Bharat, accessed March 7, 2026, https://ai4bharat.iitm.ac.in/
47. Automatic Speech Recognition - AI4Bharat, accessed March 7, 2026, https://ai4bharat.iitm.ac.in/areas/asr
48. Outbound Voice AI: 2026 Insights - Telnyx, accessed March 7, 2026, https://telnyx.com/resources/outbound-voice-ai
49. AI-Powered Customer Support: Balancing Automation and Human Touch in 2025 - LTVplus, accessed March 7, 2026, https://www.ltvplus.com/customer-service/ai-powered-customer-support/